

2 Heuristic Search

Strategy	Frontier Selection	Halts?	Space	Time
Depth-first	Last node added	No	Linear	Exp
Breadth-first	First node added	Yes	Exp	Exp
Lowest-cost-first	$\min cost(n)$	Yes	Exp	Exp
Greedy Best-first	$\min h(n)$	No	Exp	Exp
A*	$\min cost(n) + h(n)$	Yes	Exp	Exp

2 Heuristic Search-20250525140945825.webp

TL;DR

For A*:

- the heuristic $h(n)$ is an estimate of the cost from n to the goal node
- any heuristic must at least be **admissible** -> it must never overestimate the true cost
- if a heuristic is **consistent**, then $h(a) - h(b) \leq \text{edge}(a, b)$ for all a, b
 - this is the monotone condition
 - consistency -> admissibility
 - A* no longer revisits nodes
- the algorithm itself is just Dijkstra's with a heuristic

Heuristic function $h(n)$

A search heuristic $h(n)$ is an estimate of the cost of the cheapest path from node n to a goal node.

A good $h(n)$ has these properties:

1. Problem specific
2. Non-negative
3. $h(\text{goal node}) = 0$
4. $h(n)$ must be easy to compute

Lowest Cost First Search (LCFS)

Prioritize expanding the node with the lowest path cost $g(n)$ from the start.

This is just Dijkstra's algorithm. See [Dijkstra's algorithm](#).

Greedy Best First Search

BFS but the queue is a heap ordered by the heuristic $h(n)$.

```
def greedy_best_first_search(start, goal, neighbors_fn, h_fn):  
    # Priority queue with (h(n), node)  
    open_set = [(h_fn(start), start)]
```

```

came_from = {start: None}
visited = set()

while open_set:
    _, current = heapq.heappop(open_set)

    if current == goal:
        return reconstruct_path(came_from, current)

    visited.add(current)

    for neighbor in neighbors_fn(current):
        if neighbor not in visited:
            visited.add(neighbor)
            came_from[neighbor] = current
            heapq.heappush(open_set, (h_fn(neighbor), neighbor))

return None # No path found

```

Pros:

- Super easy to implement and fast if the heuristic is accurate
- Can use less memory than regular BFS or A* (not guaranteed)

Cons:

- Not optimal - if the heuristic is misleading it may pick a worse path
- Not complete - may end up in a cycle or dead end

A*

"It's just Dijkstra's but we also factor in the heuristic so that we don't explore paths that are obviously farther away from the goal."

A* is similar to Dijkstra's but when we calculate costs we also factor in a heuristic so $c(n) = g(n) + h(n)$. Where $g(n)$ is the normal Dijkstra's cost.

```

open_heap = []
g_score = default_inf()
g_score[start] = 0
heapq.heappush(open_heap, (h(start, goal), start)) # f = g + h = h at start

while open_heap:
    f_curr, current = heapq.heappop(open_heap)

    if current == goal:
        return reconstruct_path(...)

    # stale entry check: if f_curr > g_score[current] + h(current, goal): continue

    for neighbor in neighbors(current):
        tentative_g = g_score[current] + cost(current, neighbor)
        if tentative_g < g_score[neighbor]:

```

```
g_score[neighbor] = tentative_g
came_from[neighbor] = current
f_neighbor = tentative_g + h(neighbor, goal)
heapq.heappush(open_heap, (f_neighbor, neighbor))
```

Implementation notes:

- **f_cost** is used only to order the min heap
- We compare only **g_cost** scores when deciding to update path parents
- In practice, a minimum **g_cost** tracker is needed as well to detect stale heap entries (this must be in a dict and in each heap entry)
 - We may push multiple entries of the same node into the heap
 - Only the lowest cost one is up to date (valid)

How do we choose a heuristic function?

1. The heuristic MUST be admissible.
2. The heuristic SHOULD be consistent.

The heuristic function $h(n)$ is always defined as an estimate of the cheapest cost from n to the goal node.

Admissibility

The heuristic function *must be admissible* - it must never **over** estimate the cost of the cheapest path from n to the goal.

In other words: $h(n) \leq \text{true_optimal_cost}(n, \text{goal node}) = h^*(n)$

If the heuristic is admissible, then A is optimal.* (Theorem of A* Optimality)

What happens if it over estimates?

- A* thinks a good path is bad
- Explores a bad path first, ignoring a good one

Constructing an admissible heuristic

1. **Define a relaxed problem** - simplify or remove constraints on the original
 2. **Solve the relaxed problem without search**
 3. The cost of the optimal solution to the relaxed problem is an admissible heuristic for the original problem
 - **Why?** The cost of a solution with less constraints should usually be smaller than that of the original problem, which satisfies the admissibility requirement $h(n) \leq h^*(n)$
 - We should also prefer heuristics that are very different for different states
-

Some Heuristic Functions for 8-Puzzle

► Manhattan Distance Heuristic:

The sum of the Manhattan distances of the tiles from their goal positions

► Misplaced Tile Heuristic:

The number of tiles that are NOT in their goal positions

Both heuristic functions are admissible.

Initial State			Goal State		
5	3		1	2	3
8	7	6	4	5	6
2	4	1	7	8	

2 Heuristic Search-20250526175602959.webp

Dominating heuristic

Assume we have two admissible heuristics $h_a(n)$ and $h_b(n)$.

$h_a(n)$ **dominates** $h_b(n)$ if the heuristic output for $h_a(n)$ is \geq the output from $h_b(n)$ for all possible n .

If a heuristic A dominates another heuristic B, the dominating heuristic A is better.

Consistency

A heuristic is consistent if it satisfies the *monotone condition* - for any two neighbor nodes n and m :

$$h(m) - h(n) \leq \text{edge_cost}(m,n)$$

A* with multi-path pruning is only optimal if there is consistent heuristic function.

Consistency \Rightarrow admissibility

Admissibility \neq consistency (but usually, admissible heuristics are consistent)

Implications of consistency

If you find a heuristic which is consistent (and thus also admissible), then you get the following guarantees:

1. A* never needs to revisit nodes
2. A* becomes simpler and much more optimal

Cycle pruning

Check that the nodes we are trying to visit are not already on the path.

- Time complexity: Linear if using a list, $O(1)$ using a list and stack

Multipath pruning

Discard new paths to a node if we already found one.

- Use a visited set for $O(1)$ pruning
- Saves computation but increases space consumption

Can multipath pruning cause a search algorithm to fail to find the optimal solution? Say we keep the first path and prune the rest, but the first path is not the most optimal.

- **LCFS - NO.** Dijkstra's will always find the least cost path first
- **A* - YES.** The first path may not be optimal - A* with multi-path pruning is NOT optimal

Making A* work with multipath pruning

To make A* work with multipath pruning, the heuristic must be consistent (satisfies the monotone condition)

Monotone condition: for all neighbor pairs m and n ,

$$h(m) - h(n) \leq \text{edge_cost}(m,n)$$

Why this is the case:

When does multi-path pruning not work?

Assuming we have a frontier $(s \rightarrow n, \dots, s \rightarrow n')$, and we are exploring node n .

- ▶ If there exists another path through n' to n with lower f-value.
- ▶ 1) we have $h(n) + \text{cost}(n) > h(n) + \text{cost}(n') + \text{cost}(n', n)$,
e.g. $\text{cost}(n) - \text{cost}(n') > \text{cost}(n', n)$
- ▶ 2) node n is already explored, so
 $h(n) + \text{cost}(n) \leq h(n') + \text{cost}(n')$
- ▶ Combine these two, we have
 $h(n') - h(n) \geq \text{cost}(n) - \text{cost}(n') > \text{cost}(n', n)$
- ▶ Such scenario only happens when there exists two nodes n and n' with $h(n') - h(n) > \text{cost}(n', n)$.